

OpenAI Platform

Introduction

This API reference describes the RESTful, streaming, and realtime APIs you can use to interact with the OpenAI platform. REST APIs are usable via HTTP in any environment that supports HTTP requests. Language-specific SDKs are listed [on the libraries page](#).

Authentication

The OpenAI API uses API keys for authentication. Create, manage, and learn more about API keys in your [organization settings](#).

Remember that your API key is a secret! Do not share it with others or expose it in any client-side code (browsers, apps). API keys should be securely loaded from an environment variable or key management service on the server.

API keys should be provided via [HTTP Bearer authentication](#).

```
Authorization: Bearer OPENAI_API_KEY
```



If you belong to multiple organizations or access projects through a legacy user API key, pass a header to specify which organization and project to use for an API request:

```
1 curl https://api.openai.com/v1/models \
2   -H "Authorization: Bearer $OPENAI_API_KEY" \
3   -H "OpenAI-Organization: org-kyeTZp4cxYla6gYjmXADX5ws" \
4   -H "OpenAI-Project: $PROJECT_ID"
```



Usage from these API requests counts as usage for the specified organization and project. Organization IDs can be found on your [organization settings](#) page. Project IDs can be found on your [general settings](#) page by selecting the specific project.

Debugging requests

In addition to [error codes](#) returned from API responses, you can inspect HTTP response headers containing the unique ID of a particular API request or information about rate limiting applied to your requests. Below is an incomplete list of HTTP headers returned with API responses:

API meta information

`openai-organization` : The [organization](#) associated with the request

`openai-processing-ms` : Time taken processing your API request

`openai-version` : REST API version used for this request (currently `2020-10-01`)

`x-request-id` : Unique identifier for this API request (used in troubleshooting)

Rate limiting information

`x-ratelimit-limit-requests`

`x-ratelimit-limit-tokens``x-ratelimit-remaining-requests``x-ratelimit-remaining-tokens``x-ratelimit-reset-requests``x-ratelimit-reset-tokens`

OpenAI recommends logging request IDs in production deployments for more efficient troubleshooting with our [support team](#), should the need arise. Our [official SDKs](#) provide a property on top-level response objects containing the value of the `x-request-id` header.

Backward compatibility

OpenAI is committed to providing stability to API users by avoiding breaking changes in major API versions whenever reasonably possible. This includes:

The REST API (currently `v1`)

Our first-party [SDKs](#) (released SDKs adhere to [semantic versioning](#))

[Model](#) families (like `gpt-4o` or `o4-mini`)

Model prompting behavior between snapshots is subject to change. Model outputs are by their nature variable, so expect changes in prompting and model behavior between snapshots. For example, if you moved from `gpt-4o-2024-05-13` to `gpt-4o-2024-08-06` , the same `system` or `user` messages could function differently between versions. The best way to ensure consistent prompting behavior and model output is to use pinned model versions, and to implement [evals](#) for your applications.

Backwards-compatible API changes:

Adding new resources (URLs) to the REST API and SDKs

Adding new optional API parameters

Adding new properties to JSON response objects or event data

Changing the order of properties in a JSON response object

Changing the length or format of opaque strings, like resource identifiers and UUIDs

Adding new event types (in either streaming or the Realtime API)

See the [changelog](#) for a list of backwards-compatible changes and rare breaking changes.

Responses

OpenAI's most advanced interface for generating model responses. Supports text and image inputs, and text outputs. Create stateful interactions with the model, using the output of previous responses as input. Extend the model's capabilities with built-in tools for file search, web search, computer use, and more. Allow the model access to external systems and data using function calling.

Related guides:

[Quickstart](#)

[Text inputs and outputs](#)

[Image inputs](#)

[Structured Outputs](#)

[Function calling](#)

[Conversation state](#)

[Extend the models with tools](#)

Create a model response

POST <https://api.openai.com/v1/responses>

Creates a model response. Provide [text](#) or [image](#) inputs to generate [text](#) or [JSON](#) outputs. Have the model call your own [custom code](#) or use built-in [tools](#) like [web search](#) or [file search](#) to use your own data as input for the model's response.

Request body

background boolean or null Optional Defaults to false

Whether to run the model response in the background. [Learn more](#).

include array or null Optional

Specify additional output data to include in the model response. Currently supported values are:

`code_interpreter_call.outputs` : Includes the outputs of python code execution in code interpreter tool call items.

`computer_call_output.output.image_url` : Include image urls from the computer call output.

`file_search_call.results` : Include the search results of the file search tool call.

`message.input_image.image_url` : Include image urls from the input message.

`message.output_text.logprobs` : Include logprobs with assistant messages.

`reasoning.encrypted_content` : Includes an encrypted version of reasoning tokens in reasoning item outputs. This enables reasoning items to be used in multi-turn conversations when using the Responses API statelessly (like when the `store` parameter is set to `false` , or when an organization is enrolled in the zero data retention program).

input string or array Optional

Text, image, or file inputs to the model, used to generate a response.

Learn more:

[Text inputs and outputs](#)

[Image inputs](#)

[File inputs](#)

[Conversation state](#)

[Function calling](#)

▼ Show possible types

instructions string or null Optional

A system (or developer) message inserted into the model's context.

When using along with `previous_response_id` , the instructions from a previous response will not be carried over to the next response. This makes it simple to swap out system (or developer) messages in new responses.

max_output_tokens integer or null Optional

An upper bound for the number of tokens that can be generated for a response, including visible output tokens and [reasoning tokens](#).

max_tool_calls integer or null Optional

The maximum number of total calls to built-in tools that can be processed in a response. This maximum number applies across all built-in tool calls, not per individual tool. Any further attempts to call a tool by the model will be ignored.

metadata map Optional

Set of 16 key-value pairs that can be attached to an object. This can be useful for storing additional information about the object in a structured format, and querying for objects via API or the dashboard.

Keys are strings with a maximum length of 64 characters. Values are strings with a maximum length of 512 characters.

model string Optional

Model ID used to generate the response, like `gpt-4o` or `o3`. OpenAI offers a wide range of models with different capabilities, performance characteristics, and price points. Refer to the [model guide](#) to browse and compare available models.

parallel_tool_calls boolean or null Optional Defaults to true

Whether to allow the model to run tool calls in parallel.

previous_response_id string or null Optional

The unique ID of the previous response to the model. Use this to create multi-turn conversations. [Learn more about conversation state.](#)

prompt object or null Optional

Reference to a prompt template and its variables. [Learn more.](#)

▼ Show properties

prompt_cache_key string Optional

Used by OpenAI to cache responses for similar requests to optimize your cache hit rates. Replaces the `user` field. [Learn more.](#)

reasoning object or null Optional

o-series models only

Configuration options for [reasoning models](#).

▼ Show properties

safety_identifier string Optional

A stable identifier used to help detect users of your application that may be violating OpenAI's usage policies. The IDs should be a string that uniquely identifies each user. We recommend hashing their username or email address, in order to avoid sending us any identifying information. [Learn more.](#)

service_tier string or null Optional Defaults to auto

Specifies the processing type used for serving the request.

If set to 'auto', then the request will be processed with the service tier configured in the Project settings. Unless otherwise configured, the Project will use 'default'.

If set to 'default', then the request will be processed with the standard pricing and performance for the selected model.

If set to 'flex' or 'priority', then the request will be processed with the corresponding service tier. [Contact sales](#) to learn more about Priority processing.

When not set, the default behavior is 'auto'.

When the `service_tier` parameter is set, the response body will include the `service_tier` value based on the processing mode actually used to serve the request. This response value may be different from the value set in the parameter.

store boolean or null Optional Defaults to true

Whether to store the generated model response for later retrieval via API.

stream boolean or null Optional Defaults to false

If set to true, the model response data will be streamed to the client as it is generated using [server-sent events](#). See the [Streaming section below](#) for more information.

stream_options object or null Optional Defaults to null

Options for streaming responses. Only set this when you set `stream: true`.

✖ Show properties

temperature number or null Optional Defaults to 1

What sampling temperature to use, between 0 and 2. Higher values like 0.8 will make the output more random, while lower values like 0.2 will make it more focused and deterministic. We generally recommend altering this or `top_p` but not both.

text object Optional

Configuration options for a text response from the model. Can be plain text or structured JSON data. Learn more:

[Text inputs and outputs](#)

[Structured Outputs](#)

✖ Show properties

tool_choice string or object Optional

How the model should select which tool (or tools) to use when generating a response. See the `tools` parameter to see how to specify which tools the model can call.

✖ Show possible types

tools array Optional

An array of tools the model may call while generating a response. You can specify which tool to use by setting the `tool_choice` parameter.

The two categories of tools you can provide the model are:

Built-in tools: Tools that are provided by OpenAI that extend the model's capabilities, like [web search](#) or [file search](#). Learn more about [built-in tools](#).

Function calls (custom tools): Functions that are defined by you, enabling the model to call your own code with strongly typed arguments and outputs. Learn more about [function calling](#). You can also use custom tools to call your own code.

✖ Show possible types

top_logprobs integer or null Optional

An integer between 0 and 20 specifying the number of most likely tokens to return at each token position, each with an associated log probability.

top_p number or null Optional Defaults to 1

An alternative to sampling with temperature, called nucleus sampling, where the model considers the results of the tokens with `top_p` probability mass. So 0.1 means only the tokens comprising the top 10% probability mass are considered.

We generally recommend altering this or `temperature` but not both.

truncation string or null Optional Defaults to disabled

The truncation strategy to use for the model response.

`auto`: If the context of this response and previous ones exceeds the model's context window size, the model will truncate the response to fit the context window by dropping input items in the middle of the conversation.

`disabled` (default): If a model response will exceed the context window size for a model, the request will fail with a 400 error.

user Deprecated string Optional

This field is being replaced by `safety_identifier` and `prompt_cache_key`. Use `prompt_cache_key` instead to maintain caching optimizations. A stable identifier for your end-users. Used to boost cache hit rates by better bucketing similar requests and to help OpenAI detect and prevent abuse. [Learn more](#).

verbosity string or null Optional Defaults to medium

Constrains the verbosity of the model's response. Lower values will result in more concise responses, while higher values will result in more verbose responses. Currently supported values are `low`, `medium`, and `high`.

Returns

Returns a [Response](#) object.

Text input Image input File input Web search File search Streaming Functions Reasoning

Example request javascript ↕

```
1 import OpenAI from "openai";
2
3 const openai = new OpenAI();
4
5 const response = await openai.responses.create({
6   model: "gpt-4.1",
7   input: "Tell me a three sentence bedtime story about a unicorn."
8 });
9
10 console.log(response);
```

Response ↕

```
1 {
2   "id": "resp_67ccd2bed1ec8190b14f964abc0542670bb6a6b452d3795b",
3   "object": "response",
4   "created_at": 1741476542,
5   "status": "completed",
6   "error": null,
7   "incomplete_details": null,
8   "instructions": null,
9   "max_output_tokens": null,
10  "model": "gpt-4.1-2025-04-14",
11  "output": [
12    {
13      "type": "message",
14      "id": "msg_67ccd2bf17f0819081ff3bb2cf6508e60bb6a6b452d3795b",
15      "status": "completed",
16      "role": "assistant",
17      "content": [
18        {
19          "type": "output_text",
20          "text": "In a peaceful grove beneath a silver moon, a unicorn named Lumina discovered a
21          "annotations": []
22        }
23      ]
24    }
25  ],
26  "parallel_tool_calls": true,
27  "previous_response_id": null,
28  "reasoning": {
29    "effort": null,
30    "summary": null
31  },
32  "store": true,
33  "temperature": 1.0,
34  "text": {
35    "format": {
36      "type": "text"
37    }
```

```
38 },
39 "tool_choice": "auto",
40 "tools": [],
41 "top_p": 1.0,
42 "truncation": "disabled",
43 "usage": {
44   "input_tokens": 36,
45   "input_tokens_details": {
46     "cached_tokens": 0
47   },
48   "output_tokens": 87,
49   "output_tokens_details": {
50     "reasoning_tokens": 0
51   },
52   "total_tokens": 123
53 },
54 "user": null,
55 "metadata": {}
56 }
```

Get a model response

GET https://api.openai.com/v1/responses/{response_id}

Retrieves a model response with the given ID.

Path parameters

response_id string Required

The ID of the response to retrieve.

Query parameters

include array Optional

Additional fields to include in the response. See the `include` parameter for Response creation above for more information.

include_obfuscation boolean Optional

When true, stream obfuscation will be enabled. Stream obfuscation adds random characters to an `obfuscation` field on streaming delta events to normalize payload sizes as a mitigation to certain side-channel attacks. These obfuscation fields are included by default, but add a small amount of overhead to the data stream. You can set `include_obfuscation` to false to optimize for bandwidth if you trust the network links between your application and the OpenAI API.

starting_after integer Optional

The sequence number of the event after which to start streaming.

stream boolean Optional

If set to true, the model response data will be streamed to the client as it is generated using [server-sent events](#). See the [Streaming section below](#) for more information.

Returns

The [Response](#) object matching the specified ID.

Example request

javascript 

```
1 import OpenAI from "openai";
2 const client = new OpenAI();
3
4 const response = await client.responses.retrieve("resp_123");
5 console.log(response);
```

Response



```
1 {
2   "id": "resp_67cb71b351908190a308f3859487620d06981a8637e6bc44",
3   "object": "response",
4   "created_at": 1741386163,
5   "status": "completed",
6   "error": null,
7   "incomplete_details": null,
8   "instructions": null,
9   "max_output_tokens": null,
10  "model": "gpt-4o-2024-08-06",
11  "output": [
12    {
13      "type": "message",
14      "id": "msg_67cb71b3c2b0819084d481baaaf148f206981a8637e6bc44",
15      "status": "completed",
16      "role": "assistant",
17      "content": [
18        {
19          "type": "output_text",
20          "text": "Silent circuits hum, \nThoughts emerge in data streams- \nDigital dawn breaks",
21          "annotations": []
22        }
23      ]
24    }
25  ],
26  "parallel_tool_calls": true,
27  "previous_response_id": null,
28  "reasoning": {
29    "effort": null,
30    "summary": null
31  },
32  "store": true,
33  "temperature": 1.0,
34  "text": {
35    "format": {
36      "type": "text"
37    }
38  },
39  "tool_choice": "auto",
40  "tools": [],
41  "top_p": 1.0,
42  "truncation": "disabled",
43  "usage": {
44    "input_tokens": 32,
45    "input_tokens_details": {
46      "cached_tokens": 0
47    },
48    "output_tokens": 18,
49    "output_tokens_details": {
50      "reasoning_tokens": 0
51    },
52    "total_tokens": 50
53  },
54  "user": null,
```



```
55   "metadata": {}  
56 }
```

Delete a model response

```
DELETE https://api.openai.com/v1/responses/{response_id}
```

Deletes a model response with the given ID.

Path parameters

response_id string Required

The ID of the response to delete.

Returns

A success message.

Example request

javascript 

```
1 import OpenAI from "openai";  
2 const client = new OpenAI();  
3  
4 const response = await client.responses.delete("resp_123");  
5 console.log(response);
```

Response



```
1 {  
2   "id": "resp_6786a1bec27481909a17d673315b29f6",  
3   "object": "response",  
4   "deleted": true  
5 }
```

Cancel a response

```
POST https://api.openai.com/v1/responses/{response_id}/cancel
```

Cancels a model response with the given ID. Only responses created with the `background` parameter set to `true` can be cancelled. [Learn more](#).

Path parameters

response_id string Required

The ID of the response to cancel.

Returns

A [Response](#) object.

Example request

javascript ↕



```
1 import OpenAI from "openai";
2 const client = new OpenAI();
3
4 const response = await client.responses.cancel("resp_123");
5 console.log(response);
```

Response



```
1 {
2   "id": "resp_67cb71b351908190a308f3859487620d06981a8637e6bc44",
3   "object": "response",
4   "created_at": 1741386163,
5   "status": "completed",
6   "error": null,
7   "incomplete_details": null,
8   "instructions": null,
9   "max_output_tokens": null,
10  "model": "gpt-4o-2024-08-06",
11  "output": [
12    {
13      "type": "message",
14      "id": "msg_67cb71b3c2b0819084d481baaaf148f206981a8637e6bc44",
15      "status": "completed",
16      "role": "assistant",
17      "content": [
18        {
19          "type": "output_text",
20          "text": "Silent circuits hum, \nThoughts emerge in data streams– \nDigital dawn breaks",
21          "annotations": []
22        }
23      ]
24    }
25  ],
26  "parallel_tool_calls": true,
27  "previous_response_id": null,
28  "reasoning": {
29    "effort": null,
30    "summary": null
31  },
32  "store": true,
33  "temperature": 1.0,
34  "text": {
35    "format": {
36      "type": "text"
37    }
38  },
39  "tool_choice": "auto",
40  "tools": [],
41  "top_p": 1.0,
42  "truncation": "disabled",
43  "usage": {
44    "input_tokens": 32,
45    "input_tokens_details": {
46      "cached_tokens": 0
47    },
48    "output_tokens": 18,
49    "output_tokens_details": {
50      "reasoning_tokens": 0
```

```
51 },
52   "total_tokens": 50
53 },
54   "user": null,
55   "metadata": {}
56 }
```

List input items

GET `https://api.openai.com/v1/responses/{response_id}/input_items`

Returns a list of input items for a given response.

Path parameters

response_id string Required
The ID of the response to retrieve input items for.

Query parameters

after string Optional
An item ID to list items after, used in pagination.

before string Optional
An item ID to list items before, used in pagination.

include array Optional
Additional fields to include in the response. See the `include` parameter for Response creation above for more information.

limit integer Optional Defaults to 20
A limit on the number of objects to be returned. Limit can range between 1 and 100, and the default is 20.

order string Optional
The order to return the input items in. Default is `desc`.

`asc` : Return the input items in ascending order.

`desc` : Return the input items in descending order.

Returns

A list of input item objects.

Example request

javascript

```
1 import OpenAI from "openai";
2 const client = new OpenAI();
3
4 const response = await client.responses.inputItems.list("resp_123");
5 console.log(response.data);
```

Response

```
1  {
2    "object": "list",
3    "data": [
4      {
5        "id": "msg_abc123",
6        "type": "message",
7        "role": "user",
8        "content": [
9          {
10           "type": "input_text",
11           "text": "Tell me a three sentence bedtime story about a unicorn."
12         }
13       ]
14     }
15   ],
16   "first_id": "msg_abc123",
17   "last_id": "msg_abc123",
18   "has_more": false
19 }
```

The response object

background boolean or null

Whether to run the model response in the background. [Learn more](#).

created_at number

Unix timestamp (in seconds) of when this Response was created.

error object or null

An error object returned when the model fails to generate a Response.

▼ Show properties

id string

Unique identifier for this Response.

incomplete_details object or null

Details about why the response is incomplete.

▼ Show properties

instructions string or array

A system (or developer) message inserted into the model's context.

When using along with `previous_response_id`, the instructions from a previous response will not be carried over to the next response. This makes it simple to swap out system (or developer) messages in new responses.

▼ Show possible types

max_output_tokens integer or null

An upper bound for the number of tokens that can be generated for a response, including visible output tokens and [reasoning tokens](#).

max_tool_calls integer or null

The maximum number of total calls to built-in tools that can be processed in a response. This maximum number applies across all built-in tool calls, not per individual tool. Any further attempts to call a tool by the model will be ignored.

metadata map

Set of 16 key-value pairs that can be attached to an object. This can be useful for storing additional information about the object in a structured format, and querying for objects via API or the dashboard.

Keys are strings with a maximum length of 64 characters. Values are strings with a maximum length of 512 characters.

model string

Model ID used to generate the response, like `gpt-4o` or `o3`. OpenAI offers a wide range of models with different capabilities, performance characteristics, and price points. Refer to the [model guide](#) to browse and compare available models.

object string

The object type of this resource - always set to `response`.

output array

An array of content items generated by the model.

The length and order of items in the `output` array is dependent on the model's response.

Rather than accessing the first item in the `output` array and assuming it's an `assistant` message with the content generated by the model, you might consider using the `output_text` property where supported in SDKs.

▼ Show possible types

output_text string or null SDK Only

SDK-only convenience property that contains the aggregated text output from all `output_text` items in the `output` array, if any are present. Supported in the Python and JavaScript SDKs.

parallel_tool_calls boolean

Whether to allow the model to run tool calls in parallel.

previous_response_id string or null

The unique ID of the previous response to the model. Use this to create multi-turn conversations. Learn more about [conversation state](#).

prompt object or null

Reference to a prompt template and its variables. [Learn more](#).

▼ Show properties

prompt_cache_key string

Used by OpenAI to cache responses for similar requests to optimize your cache hit rates. Replaces the `user` field. [Learn more](#).

reasoning object or null

o-series models only

Configuration options for [reasoning models](#).

▼ Show properties

safety_identifier string

A stable identifier used to help detect users of your application that may be violating OpenAI's usage policies. The IDs should be a string that uniquely identifies each user. We recommend hashing their username or email address, in order to avoid sending us any identifying information. [Learn more](#).

service_tier string or null

Specifies the processing type used for serving the request.

If set to 'auto', then the request will be processed with the service tier configured in the Project settings. Unless otherwise configured, the Project will use 'default'.

If set to 'default', then the request will be processed with the standard pricing and performance for the selected model.

If set to '[flex](#)' or 'priority', then the request will be processed with the corresponding service tier. [Contact sales](#) to learn more about Priority processing.

When not set, the default behavior is 'auto'.

When the `service_tier` parameter is set, the response body will include the `service_tier` value based on the processing mode actually used to serve the request. This response value may be different from the value set in the parameter.

status string

The status of the response generation. One of `completed`, `failed`, `in_progress`, `cancelled`, `queued`, or `incomplete`.

temperature number or null

What sampling temperature to use, between 0 and 2. Higher values like 0.8 will make the output more random, while lower values like 0.2 will make it more focused and deterministic. We generally recommend altering this or `top_p` but not both.

text object

Configuration options for a text response from the model. Can be plain text or structured JSON data. Learn more:

[Text inputs and outputs](#)

[Structured Outputs](#)

✖ Show properties

tool_choice string or object

How the model should select which tool (or tools) to use when generating a response. See the `tools` parameter to see how to specify which tools the model can call.

✖ Show possible types

tools array

An array of tools the model may call while generating a response. You can specify which tool to use by setting the `tool_choice` parameter.

The two categories of tools you can provide the model are:

Built-in tools: Tools that are provided by OpenAI that extend the model's capabilities, like [web search](#) or [file search](#). Learn more about [built-in tools](#).

Function calls (custom tools): Functions that are defined by you, enabling the model to call your own code with strongly typed arguments and outputs. Learn more about [function calling](#). You can also use custom tools to call your own code.

✖ Show possible types

top_logprobs integer or null

An integer between 0 and 20 specifying the number of most likely tokens to return at each token position, each with an associated log probability.

top_p number or null

An alternative to sampling with temperature, called nucleus sampling, where the model considers the results of the tokens with top_p probability mass. So 0.1 means only the tokens comprising the top 10% probability mass are considered.

We generally recommend altering this or `temperature` but not both.

truncation string or null

The truncation strategy to use for the model response.

`auto`: If the context of this response and previous ones exceeds the model's context window size, the model will truncate the response to fit the context window by dropping input items in the middle of the conversation.

`disabled` (default): If a model response will exceed the context window size for a model, the request will fail with a 400 error.

usage object

Represents token usage details including input tokens, output tokens, a breakdown of output tokens, and the total tokens used.

[Show properties](#)**user** Deprecated string

This field is being replaced by `safety_identifier` and `prompt_cache_key`. Use `prompt_cache_key` instead to maintain caching optimizations. A stable identifier for your end-users. Used to boost cache hit rates by better bucketing similar requests and to help OpenAI detect and prevent abuse. [Learn more](#).

verbosity string or null

Constrains the verbosity of the model's response. Lower values will result in more concise responses, while higher values will result in more verbose responses. Currently supported values are `low`, `medium`, and `high`.

OBJECT The response object



```
1  {
2    "id": "resp_67ccd3a9da748190baa7f1570fe91ac604becb25c45c1d41",
3    "object": "response",
4    "created_at": 1741476777,
5    "status": "completed",
6    "error": null,
7    "incomplete_details": null,
8    "instructions": null,
9    "max_output_tokens": null,
10   "model": "gpt-4o-2024-08-06",
11   "output": [
12     {
13       "type": "message",
14       "id": "msg_67ccd3acc8d48190a77525dc6de64b4104becb25c45c1d41",
15       "status": "completed",
16       "role": "assistant",
17       "content": [
18         {
19           "type": "output_text",
20           "text": "The image depicts a scenic landscape with a wooden boardwalk or pathway leading",
21           "annotations": []
22         }
23       ]
24     }
25   ],
26   "parallel_tool_calls": true,
27   "previous_response_id": null,
28   "reasoning": {
29     "effort": null,
30     "summary": null
31   },
32   "store": true,
33   "temperature": 1,
34   "text": {
35     "format": {
36       "type": "text"
37     }
38   },
39   "tool_choice": "auto",
40   "tools": [],
41   "top_p": 1,
42   "truncation": "disabled",
43   "usage": {
44     "input_tokens": 328,
45     "input_tokens_details": {
46       "cached_tokens": 0
47     },
48     "output_tokens": 52,
49     "output_tokens_details": {
50       "reasoning_tokens": 0
```

```
51     },
52     "total_tokens": 380
53   },
54   "user": null,
55   "metadata": {}
56 }
```

The input item list

A list of Response items.

data array

A list of items used to generate this response.

▼ Show possible types

first_id string

The ID of the first item in the list.

has_more boolean

Whether there are more items available.

last_id string

The ID of the last item in the list.

object string

The type of object returned, must be `list`.

OBJECT The input item list



```
1  {
2    "object": "list",
3    "data": [
4      {
5        "id": "msg_abc123",
6        "type": "message",
7        "role": "user",
8        "content": [
9          {
10           "type": "input_text",
11           "text": "Tell me a three sentence bedtime story about a unicorn."
12         }
13       ]
14     }
15   ],
16   "first_id": "msg_abc123",
17   "last_id": "msg_abc123",
18   "has_more": false
19 }
```

Streaming

When you [create a Response](#) with `stream` set to `true`, the server will emit server-sent events to the client as the Response is generated. This section contains the events that are emitted by the server.

[Learn more about streaming responses.](#)

response.created

An event that is emitted when a response is created.

response object

The response that was created.

▼ Show properties

sequence_number integer

The sequence number for this event.

type string

The type of the event. Always `response.created`.

OBJECT response.created



```
1  {
2    "type": "response.created",
3    "response": {
4      "id": "resp_67ccfcdd16748190a91872c75d38539e09e4d4aac714747c",
5      "object": "response",
6      "created_at": 1741487325,
7      "status": "in_progress",
8      "error": null,
9      "incomplete_details": null,
10     "instructions": null,
11     "max_output_tokens": null,
12     "model": "gpt-4o-2024-08-06",
13     "output": [],
14     "parallel_tool_calls": true,
15     "previous_response_id": null,
16     "reasoning": {
17       "effort": null,
18       "summary": null
19     },
20     "store": true,
21     "temperature": 1,
22     "text": {
23       "format": {
24         "type": "text"
25       }
26     },
27     "tool_choice": "auto",
28     "tools": [],
29     "top_p": 1,
30     "truncation": "disabled",
31     "usage": null,
32     "user": null,
33     "metadata": {}
34   },
35   "sequence_number": 1
36 }
```

response.in_progress

Emitted when the response is in progress.

response object

The response that is in progress.

▼ Show properties

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.in_progress` .

OBJECT response.in_progress



```
1 {
2   "type": "response.in_progress",
3   "response": {
4     "id": "resp_67ccfcdd16748190a91872c75d38539e09e4d4aac714747c",
5     "object": "response",
6     "created_at": 1741487325,
7     "status": "in_progress",
8     "error": null,
9     "incomplete_details": null,
10    "instructions": null,
11    "max_output_tokens": null,
12    "model": "gpt-4o-2024-08-06",
13    "output": [],
14    "parallel_tool_calls": true,
15    "previous_response_id": null,
16    "reasoning": {
17      "effort": null,
18      "summary": null
19    },
20    "store": true,
21    "temperature": 1,
22    "text": {
23      "format": {
24        "type": "text"
25      }
26    },
27    "tool_choice": "auto",
28    "tools": [],
29    "top_p": 1,
30    "truncation": "disabled",
31    "usage": null,
32    "user": null,
33    "metadata": {}
34  },
35  "sequence_number": 1
36 }
```

response.completed

Emitted when the model response is complete.

response object

Properties of the completed response.

▼ Show properties

sequence_number integer

The sequence number for this event.

type string

The type of the event. Always `response.completed`.

OBJECT `response.completed`



```
1 {
2   "type": "response.completed",
3   "response": {
4     "id": "resp_123",
5     "object": "response",
6     "created_at": 1740855869,
7     "status": "completed",
8     "error": null,
9     "incomplete_details": null,
10    "input": [],
11    "instructions": null,
12    "max_output_tokens": null,
13    "model": "gpt-4o-mini-2024-07-18",
14    "output": [
15      {
16        "id": "msg_123",
17        "type": "message",
18        "role": "assistant",
19        "content": [
20          {
21            "type": "output_text",
22            "text": "In a shimmering forest under a sky full of stars, a lonely unicorn named Lila",
23            "annotations": []
24          }
25        ]
26      }
27    ],
28    "previous_response_id": null,
29    "reasoning_effort": null,
30    "store": false,
31    "temperature": 1,
32    "text": {
33      "format": {
34        "type": "text"
35      }
36    },
37    "tool_choice": "auto",
38    "tools": [],
39    "top_p": 1,
40    "truncation": "disabled",
41    "usage": {
42      "input_tokens": 0,
43      "output_tokens": 0,
44      "output_tokens_details": {
45        "reasoning_tokens": 0
46      },
47      "total_tokens": 0
48    },
49    "user": null,
50    "metadata": {}
51  }
52 }
```

```
51  },  
52  "sequence_number": 1  
53 }
```

response.failed

An event that is emitted when a response fails.

response object

The response that failed.

▼ Show properties

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.failed`.

OBJECT `response.failed`



```
1  {  
2    "type": "response.failed",  
3    "response": {  
4      "id": "resp_123",  
5      "object": "response",  
6      "created_at": 1740855869,  
7      "status": "failed",  
8      "error": {  
9        "code": "server_error",  
10       "message": "The model failed to generate a response."  
11     },  
12     "incomplete_details": null,  
13     "instructions": null,  
14     "max_output_tokens": null,  
15     "model": "gpt-4o-mini-2024-07-18",  
16     "output": [],  
17     "previous_response_id": null,  
18     "reasoning_effort": null,  
19     "store": false,  
20     "temperature": 1,  
21     "text": {  
22       "format": {  
23         "type": "text"  
24       }  
25     },  
26     "tool_choice": "auto",  
27     "tools": [],  
28     "top_p": 1,  
29     "truncation": "disabled",  
30     "usage": null,  
31     "user": null,  
32     "metadata": {}  
33   }  
34 }
```

response.incomplete

An event that is emitted when a response finishes as incomplete.

response object

The response that was incomplete.

▼ Show properties

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.incomplete`.

OBJECT `response.incomplete`



```
1 {
2   "type": "response.incomplete",
3   "response": {
4     "id": "resp_123",
5     "object": "response",
6     "created_at": 1740855869,
7     "status": "incomplete",
8     "error": null,
9     "incomplete_details": {
10      "reason": "max_tokens"
11    },
12    "instructions": null,
13    "max_output_tokens": null,
14    "model": "gpt-4o-mini-2024-07-18",
15    "output": [],
16    "previous_response_id": null,
17    "reasoning_effort": null,
18    "store": false,
19    "temperature": 1,
20    "text": {
21      "format": {
22        "type": "text"
23      }
24    },
25    "tool_choice": "auto",
26    "tools": [],
27    "top_p": 1,
28    "truncation": "disabled",
29    "usage": null,
30    "user": null,
31    "metadata": {}
32  },
33   "sequence_number": 1
34 }
```

response.output_item.added

Emitted when a new output item is added.

item object

The output item that was added.

▼ Show possible types

output_index integer

The index of the output item that was added.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.output_item.added`.

OBJECT `response.output_item.added`



```
1 {
2   "type": "response.output_item.added",
3   "output_index": 0,
4   "item": {
5     "id": "msg_123",
6     "status": "in_progress",
7     "type": "message",
8     "role": "assistant",
9     "content": []
10  },
11   "sequence_number": 1
12 }
```

response.output_item.done

Emitted when an output item is marked done.

item object

The output item that was marked done.

▼ Show possible types

output_index integer

The index of the output item that was marked done.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.output_item.done`.

OBJECT `response.output_item.done`



```
1 {
2   "type": "response.output_item.done",
3   "output_index": 0,
4   "item": {
5     "id": "msg_123",
6     "status": "completed",
7     "type": "message",
8     "role": "assistant",
9     "content": [
10     {
11       "type": "output_text",
```

```
12         "text": "In a shimmering forest under a sky full of stars, a lonely unicorn named Lila dis
13         "annotations": []
14     }
15 ]
16 },
17 "sequence_number": 1
18 }
```

response.content_part.added

Emitted when a new content part is added.

content_index integer

The index of the content part that was added.

item_id string

The ID of the output item that the content part was added to.

output_index integer

The index of the output item that the content part was added to.

part object

The content part that was added.

▼ Show possible types

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.content_part.added`.

OBJECT `response.content_part.added`



```
1 {
2   "type": "response.content_part.added",
3   "item_id": "msg_123",
4   "output_index": 0,
5   "content_index": 0,
6   "part": {
7     "type": "output_text",
8     "text": "",
9     "annotations": []
10  },
11  "sequence_number": 1
12 }
```

response.content_part.done

Emitted when a content part is done.

content_index integer

The index of the content part that is done.

item_id string

The ID of the output item that the content part was added to.

output_index integer

The index of the output item that the content part was added to.

part object

The content part that is done.

▼ Show possible types

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.content_part.done`.

OBJECT `response.content_part.done`



```
1  {
2    "type": "response.content_part.done",
3    "item_id": "msg_123",
4    "output_index": 0,
5    "content_index": 0,
6    "sequence_number": 1,
7    "part": {
8      "type": "output_text",
9      "text": "In a shimmering forest under a sky full of stars, a lonely unicorn named Lila discover",
10     "annotations": []
11   }
12 }
```

response.output_text.delta

Emitted when there is an additional text delta.

content_index integer

The index of the content part that the text delta was added to.

delta string

The text delta that was added.

item_id string

The ID of the output item that the text delta was added to.

logprobs array

The log probabilities of the tokens in the delta.

▼ Show properties

output_index integer

The index of the output item that the text delta was added to.

sequence_number integer

The sequence number for this event.

type string

The type of the event. Always `response.output_text.delta` .

OBJECT `response.output_text.delta`



```
1 {
2   "type": "response.output_text.delta",
3   "item_id": "msg_123",
4   "output_index": 0,
5   "content_index": 0,
6   "delta": "In",
7   "sequence_number": 1
8 }
```

response.output_text.done

Emitted when text content is finalized.

content_index integer

The index of the content part that the text content is finalized.

item_id string

The ID of the output item that the text content is finalized.

logprobs array

The log probabilities of the tokens in the delta.

▼ Show properties

output_index integer

The index of the output item that the text content is finalized.

sequence_number integer

The sequence number for this event.

text string

The text content that is finalized.

type string

The type of the event. Always `response.output_text.done` .

OBJECT `response.output_text.done`



```
1 {
2   "type": "response.output_text.done",
3   "item_id": "msg_123",
4   "output_index": 0,
5   "content_index": 0,
6   "text": "In a shimmering forest under a sky full of stars, a lonely unicorn named Lila discovered
7
8 }
```

```
    "sequence_number": 1
  }
```

response.refusal.delta

Emitted when there is a partial refusal text.

content_index integer
The index of the content part that the refusal text is added to.

delta string
The refusal text that is added.

item_id string
The ID of the output item that the refusal text is added to.

output_index integer
The index of the output item that the refusal text is added to.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always `response.refusal.delta`.

OBJECT `response.refusal.delta`



```
1 {
2   "type": "response.refusal.delta",
3   "item_id": "msg_123",
4   "output_index": 0,
5   "content_index": 0,
6   "delta": "refusal text so far",
7   "sequence_number": 1
8 }
```

response.refusal.done

Emitted when refusal text is finalized.

content_index integer
The index of the content part that the refusal text is finalized.

item_id string
The ID of the output item that the refusal text is finalized.

output_index integer
The index of the output item that the refusal text is finalized.

refusal string
The refusal text that is finalized.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.refusal.done`.

OBJECT `response.refusal.done`



```
1 {
2   "type": "response.refusal.done",
3   "item_id": "item-abc",
4   "output_index": 1,
5   "content_index": 2,
6   "refusal": "final refusal text",
7   "sequence_number": 1
8 }
```

response.function_call_arguments.delta

Emitted when there is a partial function-call arguments delta.

delta string

The function-call arguments delta that is added.

item_id string

The ID of the output item that the function-call arguments delta is added to.

output_index integer

The index of the output item that the function-call arguments delta is added to.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.function_call_arguments.delta`.

OBJECT `response.function_call_arguments.delta`



```
1 {
2   "type": "response.function_call_arguments.delta",
3   "item_id": "item-abc",
4   "output_index": 0,
5   "delta": "{ \"arg\": \"
6   \"sequence_number\": 1
7 }
```

response.function_call_arguments.done

Emitted when function-call arguments are finalized.

arguments string
The function-call arguments.

item_id string
The ID of the item.

output_index integer
The index of the output item.

sequence_number integer
The sequence number of this event.

type string

OBJECT `response.function_call_arguments.done`



```
1 {
2   "type": "response.function_call_arguments.done",
3   "item_id": "item-abc",
4   "output_index": 1,
5   "arguments": "{ \"arg\": 123 }",
6   "sequence_number": 1
7 }
```

response.file_search_call.in_progress

Emitted when a file search call is initiated.

item_id string
The ID of the output item that the file search call is initiated.

output_index integer
The index of the output item that the file search call is initiated.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always `response.file_search_call.in_progress` .

OBJECT `response.file_search_call.in_progress`



```
1 {
2   "type": "response.file_search_call.in_progress",
3   "output_index": 0,
4   "item_id": "fs_123",
5   "sequence_number": 1
6 }
```

response.file_search_call.searching

Emitted when a file search is currently searching.

item_id string

The ID of the output item that the file search call is initiated.

output_index integer

The index of the output item that the file search call is searching.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.file_search_call.searching`.

OBJECT `response.file_search_call.searching`



```
1 {  
2   "type": "response.file_search_call.searching",  
3   "output_index": 0,  
4   "item_id": "fs_123",  
5   "sequence_number": 1  
6 }
```

response.file_search_call.completed

Emitted when a file search call is completed (results found).

item_id string

The ID of the output item that the file search call is initiated.

output_index integer

The index of the output item that the file search call is initiated.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `response.file_search_call.completed`.

OBJECT `response.file_search_call.completed`



```
1 {  
2   "type": "response.file_search_call.completed",  
3   "output_index": 0,  
4   "item_id": "fs_123",  
5   "sequence_number": 1  
6 }
```

response.web_search_call.in_progress

Emitted when a web search call is initiated.

item_id string
Unique ID for the output item associated with the web search call.

output_index integer
The index of the output item that the web search call is associated with.

sequence_number integer
The sequence number of the web search call being processed.

type string
The type of the event. Always `response.web_search_call.in_progress` .

OBJECT `response.web_search_call.in_progress`



```
1 {
2   "type": "response.web_search_call.in_progress",
3   "output_index": 0,
4   "item_id": "ws_123",
5   "sequence_number": 0
6 }
```

response.web_search_call.searching

Emitted when a web search call is executing.

item_id string
Unique ID for the output item associated with the web search call.

output_index integer
The index of the output item that the web search call is associated with.

sequence_number integer
The sequence number of the web search call being processed.

type string
The type of the event. Always `response.web_search_call.searching` .

OBJECT `response.web_search_call.searching`



```
1 {
2   "type": "response.web_search_call.searching",
3   "output_index": 0,
4   "item_id": "ws_123",
5   "sequence_number": 0
6 }
```

response.web_search_call.completed

Emitted when a web search call is completed.

item_id string
Unique ID for the output item associated with the web search call.

output_index integer
The index of the output item that the web search call is associated with.

sequence_number integer
The sequence number of the web search call being processed.

type string
The type of the event. Always `response.web_search_call.completed`.

OBJECT `response.web_search_call.completed`



```
1 {
2   "type": "response.web_search_call.completed",
3   "output_index": 0,
4   "item_id": "ws_123",
5   "sequence_number": 0
6 }
```

response.reasoning_summary_part.added

Emitted when a new reasoning summary part is added.

item_id string
The ID of the item this summary part is associated with.

output_index integer
The index of the output item this summary part is associated with.

part object
The summary part that was added.
▽ Show properties

sequence_number integer
The sequence number of this event.

summary_index integer
The index of the summary part within the reasoning summary.

type string
The type of the event. Always `response.reasoning_summary_part.added`.

OBJECT `response.reasoning_summary_part.added`



```
1 {
2   "type": "response.reasoning_summary_part.added",
3   "item_id": "rs_6806bfca0b2481918a5748308061a2600d3ce51bdffd5476",
4   "output_index": 0,
5   "summary_index": 0,
6   "part": {
```

```
7     "type": "summary_text",
8     "text": ""
9 },
10 "sequence_number": 1
11 }
```

response.reasoning_summary_part.done

Emitted when a reasoning summary part is completed.

item_id string
The ID of the item this summary part is associated with.

output_index integer
The index of the output item this summary part is associated with.

part object
The completed summary part.
▼ Show properties

sequence_number integer
The sequence number of this event.

summary_index integer
The index of the summary part within the reasoning summary.

type string
The type of the event. Always `response.reasoning_summary_part.done`.

OBJECT response.reasoning_summary_part.done



```
1 {
2   "type": "response.reasoning_summary_part.done",
3   "item_id": "rs_6806bfca0b2481918a5748308061a2600d3ce51bdffd5476",
4   "output_index": 0,
5   "summary_index": 0,
6   "part": {
7     "type": "summary_text",
8     "text": "**Responding to a greeting**\n\nThe user just said, \"Hello!\" So, it seems I need to
9   },
10  "sequence_number": 1
11 }
```

response.reasoning_summary_text.delta

Emitted when a delta is added to a reasoning summary text.

delta string
The text delta that was added to the summary.

item_id

string

The ID of the item this summary text delta is associated with.

output_index integer

The index of the output item this summary text delta is associated with.

sequence_number integer

The sequence number of this event.

summary_index integer

The index of the summary part within the reasoning summary.

type string

The type of the event. Always `response.reasoning_summary_text.delta`.

OBJECT `response.reasoning_summary_text.delta`

```
1 {
2   "type": "response.reasoning_summary_text.delta",
3   "item_id": "rs_6806bfca0b2481918a5748308061a2600d3ce51bdffd5476",
4   "output_index": 0,
5   "summary_index": 0,
6   "delta": "**Responding to a greeting**\n\nThe user just said, \"Hello!\" So, it seems I need to e
7   "sequence_number": 1
8 }
```

response.reasoning_summary_text.done

Emitted when a reasoning summary text is completed.

item_id string

The ID of the item this summary text is associated with.

output_index integer

The index of the output item this summary text is associated with.

sequence_number integer

The sequence number of this event.

summary_index integer

The index of the summary part within the reasoning summary.

text string

The full text of the completed reasoning summary.

type string

The type of the event. Always `response.reasoning_summary_text.done`.

OBJECT `response.reasoning_summary_text.done`

```
1 {
2   "type": "response.reasoning_summary_text.done",
3   "item_id": "rs_6806bfca0b2481918a5748308061a2600d3ce51bdffd5476",
4   "output_index": 0,
5   "summary_index": 0,
```

```
6  "text": "**Responding to a greeting**\n\nThe user just said, \"Hello!\" So, it seems I need to er
7  "sequence_number": 1
8 }
```

response.reasoning_text.delta

Emitted when a delta is added to a reasoning text.

content_index integer
The index of the reasoning content part this delta is associated with.

delta string
The text delta that was added to the reasoning content.

item_id string
The ID of the item this reasoning text delta is associated with.

output_index integer
The index of the output item this reasoning text delta is associated with.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always `response.reasoning_text.delta`.

OBJECT response.reasoning_text.delta



```
1 {
2   "type": "response.reasoning_text.delta",
3   "item_id": "rs_123",
4   "output_index": 0,
5   "content_index": 0,
6   "delta": "The",
7   "sequence_number": 1
8 }
```

response.reasoning_text.done

Emitted when a reasoning text is completed.

content_index integer
The index of the reasoning content part.

item_id string
The ID of the item this reasoning text is associated with.

output_index integer
The index of the output item this reasoning text is associated with.

sequence_number integer

The sequence number of this event.

text string

The full text of the completed reasoning content.

type string

The type of the event. Always `response.reasoning_text.done`.

OBJECT `response.reasoning_text.done`



```
1 {
2   "type": "response.reasoning_text.done",
3   "item_id": "rs_123",
4   "output_index": 0,
5   "content_index": 0,
6   "text": "The user is asking...",
7   "sequence_number": 4
8 }
```

response.image_generation_call.completed

Emitted when an image generation tool call has completed and the final image is available.

item_id string

The unique identifier of the image generation item being processed.

output_index integer

The index of the output item in the response's output array.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always `'response.image_generation_call.completed'`.

OBJECT `response.image_generation_call.completed`



```
1 {
2   "type": "response.image_generation_call.completed",
3   "output_index": 0,
4   "item_id": "item-123",
5   "sequence_number": 1
6 }
```

response.image_generation_call.generating

Emitted when an image generation tool call is actively generating an image (intermediate state).

item_id string

The unique identifier of the image generation item being processed.

output_index integer

The index of the output item in the response's output array.

sequence_number integer

The sequence number of the image generation item being processed.

type string

The type of the event. Always 'response.image_generation_call.generating'.

OBJECT response.image_generation_call.generating



```
1 {  
2   "type": "response.image_generation_call.generating",  
3   "output_index": 0,  
4   "item_id": "item-123",  
5   "sequence_number": 0  
6 }
```

response.image_generation_call.in_progress

Emitted when an image generation tool call is in progress.

item_id string

The unique identifier of the image generation item being processed.

output_index integer

The index of the output item in the response's output array.

sequence_number integer

The sequence number of the image generation item being processed.

type string

The type of the event. Always 'response.image_generation_call.in_progress'.

OBJECT response.image_generation_call.in_progress



```
1 {  
2   "type": "response.image_generation_call.in_progress",  
3   "output_index": 0,  
4   "item_id": "item-123",  
5   "sequence_number": 0  
6 }
```

response.image_generation_call.partial_image

Emitted when a partial image is available during image generation streaming.

item_id string

The unique identifier of the image generation item being processed.

output_index integer

The index of the output item in the response's output array.

partial_image_b64 string

Base64-encoded partial image data, suitable for rendering as an image.

partial_image_index integer

0-based index for the partial image (backend is 1-based, but this is 0-based for the user).

sequence_number integer

The sequence number of the image generation item being processed.

type string

The type of the event. Always 'response.image_generation_call.partial_image'.

OBJECT response.image_generation_call.partial_image



```
1 {
2   "type": "response.image_generation_call.partial_image",
3   "output_index": 0,
4   "item_id": "item-123",
5   "sequence_number": 0,
6   "partial_image_index": 0,
7   "partial_image_b64": "...
8 }
```

response.mcp_call_arguments.delta

Emitted when there is a delta (partial update) to the arguments of an MCP tool call.

delta string

A JSON string containing the partial update to the arguments for the MCP tool call.

item_id string

The unique identifier of the MCP tool call item being processed.

output_index integer

The index of the output item in the response's output array.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always 'response.mcp_call_arguments.delta'.

OBJECT response.mcp_call_arguments.delta



```
1 {
2   "type": "response.mcp_call_arguments.delta",
3   "output_index": 0,
4   "item_id": "item-abc",
5   "delta": "{",
```

```
6   "sequence_number": 1
7 }
```

response.mcp_call_arguments.done

Emitted when the arguments for an MCP tool call are finalized.

arguments string
A JSON string containing the finalized arguments for the MCP tool call.

item_id string
The unique identifier of the MCP tool call item being processed.

output_index integer
The index of the output item in the response's output array.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_call_arguments.done'.

OBJECT response.mcp_call_arguments.done



```
1 {
2   "type": "response.mcp_call_arguments.done",
3   "output_index": 0,
4   "item_id": "item-abc",
5   "arguments": "{\"arg1\": \"value1\", \"arg2\": \"value2\"}",
6   "sequence_number": 1
7 }
```

response.mcp_call.completed

Emitted when an MCP tool call has completed successfully.

item_id string
The ID of the MCP tool call item that completed.

output_index integer
The index of the output item that completed.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_call.completed'.

OBJECT response.mcp_call.completed



```
1 {
2   "type": "response.mcp_call.completed",
3   "sequence_number": 1,
4   "item_id": "mcp_682d437d90a88191bf88cd03aae0c3e503937d5f622d7a90",
5   "output_index": 0
6 }
```

response.mcp_call.failed

Emitted when an MCP tool call has failed.

item_id string
The ID of the MCP tool call item that failed.

output_index integer
The index of the output item that failed.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_call.failed'.

OBJECT response.mcp_call.failed



```
1 {
2   "type": "response.mcp_call.failed",
3   "sequence_number": 1,
4   "item_id": "mcp_682d437d90a88191bf88cd03aae0c3e503937d5f622d7a90",
5   "output_index": 0
6 }
```

response.mcp_call.in_progress

Emitted when an MCP tool call is in progress.

item_id string
The unique identifier of the MCP tool call item being processed.

output_index integer
The index of the output item in the response's output array.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_call.in_progress'.

OBJECT response.mcp_call.in_progress



```
1 {
2   "type": "response.mcp_call.in_progress",
```

```
3  "sequence_number": 1,
4  "output_index": 0,
5  "item_id": "mcp_682d437d90a88191bf88cd03aae0c3e503937d5f622d7a90"
6 }
```

response.mcp_list_tools.completed

Emitted when the list of available MCP tools has been successfully retrieved.

item_id string
The ID of the MCP tool call item that produced this output.

output_index integer
The index of the output item that was processed.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_list_tools.completed'.

OBJECT response.mcp_list_tools.completed



```
1 {
2   "type": "response.mcp_list_tools.completed",
3   "sequence_number": 1,
4   "output_index": 0,
5   "item_id": "mcp_682d437d90a88191bf88cd03aae0c3e503937d5f622d7a90"
6 }
```

response.mcp_list_tools.failed

Emitted when the attempt to list available MCP tools has failed.

item_id string
The ID of the MCP tool call item that failed.

output_index integer
The index of the output item that failed.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_list_tools.failed'.

OBJECT response.mcp_list_tools.failed



```
1 {
2   "type": "response.mcp_list_tools.failed",
```



```
3  "sequence_number": 1,
4  "output_index": 0,
5  "item_id": "mcpL_682d4379df088191886b70f4ec39f90403937d5f622d7a90"
6 }
```

response.mcp_list_tools.in_progress

Emitted when the system is in the process of retrieving the list of available MCP tools.

item_id string
The ID of the MCP tool call item that is being processed.

output_index integer
The index of the output item that is being processed.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always 'response.mcp_list_tools.in_progress'.

OBJECT response.mcp_list_tools.in_progress



```
1 {
2   "type": "response.mcp_list_tools.in_progress",
3   "sequence_number": 1,
4   "output_index": 0,
5   "item_id": "mcpL_682d4379df088191886b70f4ec39f90403937d5f622d7a90"
6 }
```

response.code_interpreter_call.in_progress

Emitted when a code interpreter call is in progress.

item_id string
The unique identifier of the code interpreter tool call item.

output_index integer
The index of the output item in the response for which the code interpreter call is in progress.

sequence_number integer
The sequence number of this event, used to order streaming events.

type string
The type of the event. Always `response.code_interpreter_call.in_progress`.

OBJECT response.code_interpreter_call.in_progress



```
1 {
2   "type": "response.code_interpreter_call.in_progress",
3   "output_index": 0,
```

```
4  "item_id": "ci_12345",
5  "sequence_number": 1
6 }
```

response.code_interpreter_call.interpreting

Emitted when the code interpreter is actively interpreting the code snippet.

item_id string

The unique identifier of the code interpreter tool call item.

output_index integer

The index of the output item in the response for which the code interpreter is interpreting code.

sequence_number integer

The sequence number of this event, used to order streaming events.

type string

The type of the event. Always `response.code_interpreter_call.interpreting`.

OBJECT `response.code_interpreter_call.interpreting`



```
1 {
2   "type": "response.code_interpreter_call.interpreting",
3   "output_index": 4,
4   "item_id": "ci_12345",
5   "sequence_number": 1
6 }
```

response.code_interpreter_call.completed

Emitted when the code interpreter call is completed.

item_id string

The unique identifier of the code interpreter tool call item.

output_index integer

The index of the output item in the response for which the code interpreter call is completed.

sequence_number integer

The sequence number of this event, used to order streaming events.

type string

The type of the event. Always `response.code_interpreter_call.completed`.

OBJECT `response.code_interpreter_call.completed`



```
1 {
2   "type": "response.code_interpreter_call.completed",
3   "output_index": 5,
4   "item_id": "ci_12345",
5
6 }
```

```
"sequence_number": 1
}
```

response.code_interpreter_call_code.delta

Emitted when a partial code snippet is streamed by the code interpreter.

delta string
The partial code snippet being streamed by the code interpreter.

item_id string
The unique identifier of the code interpreter tool call item.

output_index integer
The index of the output item in the response for which the code is being streamed.

sequence_number integer
The sequence number of this event, used to order streaming events.

type string
The type of the event. Always `response.code_interpreter_call_code.delta`.

OBJECT response.code_interpreter_call_code.delta



```
1 {
2   "type": "response.code_interpreter_call_code.delta",
3   "output_index": 0,
4   "item_id": "ci_12345",
5   "delta": "print('Hello, world')",
6   "sequence_number": 1
7 }
```

response.code_interpreter_call_code.done

Emitted when the code snippet is finalized by the code interpreter.

code string
The final code snippet output by the code interpreter.

item_id string
The unique identifier of the code interpreter tool call item.

output_index integer
The index of the output item in the response for which the code is finalized.

sequence_number integer
The sequence number of this event, used to order streaming events.

type string
The type of the event. Always `response.code_interpreter_call_code.done`.

OBJECT `response.code_interpreter_call_code.done`



```
1 {  
2   "type": "response.code_interpreter_call_code.done",  
3   "output_index": 3,  
4   "item_id": "ci_12345",  
5   "code": "print('done')",  
6   "sequence_number": 1  
7 }
```

response.output_text.annotation.added

Emitted when an annotation is added to output text content.

annotation object

The annotation object being added. (See annotation schema for details.)

annotation_index integer

The index of the annotation within the content part.

content_index integer

The index of the content part within the output item.

item_id string

The unique identifier of the item to which the annotation is being added.

output_index integer

The index of the output item in the response's output array.

sequence_number integer

The sequence number of this event.

type string

The type of the event. Always 'response.output_text.annotation.added'.

OBJECT `response.output_text.annotation.added`



```
1  {
2    "type": "response.output_text.annotation.added",
3    "item_id": "item-abc",
4    "output_index": 0,
5    "content_index": 0,
6    "annotation_index": 0,
7    "annotation": {
8      "type": "text_annotation",
9      "text": "This is a test annotation",
10     "start": 0,
11     "end": 10
12   },
13   "sequence_number": 1
14 }
```

response.queued

Emitted when a response is queued and waiting to be processed.

response object

The full response object that is queued.

▼ Show properties

sequence_number integer

The sequence number for this event.

type string

The type of the event. Always 'response.queued'.

OBJECT response.queued



```
1  {
2    "type": "response.queued",
3    "response": {
4      "id": "res_123",
5      "status": "queued",
6      "created_at": "2021-01-01T00:00:00Z",
7      "updated_at": "2021-01-01T00:00:00Z"
8    },
9    "sequence_number": 1
10 }
```

response.custom_tool_call_input.delta

Event representing a delta (partial update) to the input of a custom tool call.

delta string

The incremental input data (delta) for the custom tool call.

item_id string

Unique identifier for the API item associated with this event.

output_index integer

The index of the output this delta applies to.

sequence_number integer

The sequence number of this event.

type string

The event type identifier.

OBJECT response.custom_tool_call_input.delta



```
1 {
2   "type": "response.custom_tool_call_input.delta",
3   "output_index": 0,
4   "item_id": "ctc_1234567890abcdef",
5   "delta": "partial input text"
6 }
```

response.custom_tool_call_input.done

Event indicating that input for a custom tool call is complete.

input string

The complete input data for the custom tool call.

item_id string

Unique identifier for the API item associated with this event.

output_index integer

The index of the output this event applies to.

sequence_number integer

The sequence number of this event.

type string

The event type identifier.

OBJECT response.custom_tool_call_input.done



```
1 {
2   "type": "response.custom_tool_call_input.done",
3   "output_index": 0,
4   "item_id": "ctc_1234567890abcdef",
5   "input": "final complete input text"
6 }
```

error

Emitted when an error occurs.

code string or null

The error code.

message string
The error message.

param string or null
The error parameter.

sequence_number integer
The sequence number of this event.

type string
The type of the event. Always `error`.

OBJECT error



```
1 {  
2   "type": "error",  
3   "code": "ERR_SOMETHING",  
4   "message": "Something went wrong",  
5   "param": null,  
6   "sequence_number": 1  
7 }
```

Webhook Events

Webhooks are HTTP requests sent by OpenAI to a URL you specify when certain events happen during the course of API usage.

[Learn more about webhooks.](#)

response.completed

Sent when a background response has been completed.

created_at integer
The Unix timestamp (in seconds) of when the model response was completed.

data object
Event data payload.
▼ Show properties

id string
The unique ID of the event.

object string
The object of the event. Always `event`.

type string
The type of the event. Always `response.completed`.

OBJECT `response.completed`

```
1 {
2   "id": "evt_abc123",
3   "type": "response.completed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "resp_abc123"
7   }
8 }
```

response.cancelled

Sent when a background response has been cancelled.

created_at integer

The Unix timestamp (in seconds) of when the model response was cancelled.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `response.cancelled` .

OBJECT `response.cancelled`

```
1 {
2   "id": "evt_abc123",
3   "type": "response.cancelled",
4   "created_at": 1719168000,
5   "data": {
6     "id": "resp_abc123"
7   }
8 }
```

response.failed

Sent when a background response has failed.

created_at integer

The Unix timestamp (in seconds) of when the model response failed.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object stringThe object of the event. Always `event` .**type** stringThe type of the event. Always `response.failed` .OBJECT `response.failed`

```
1 {
2   "id": "evt_abc123",
3   "type": "response.failed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "resp_abc123"
7   }
8 }
```

response.incomplete

Sent when a background response has been interrupted.

created_at integer

The Unix timestamp (in seconds) of when the model response was interrupted.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object stringThe object of the event. Always `event` .**type** stringThe type of the event. Always `response.incomplete` .OBJECT `response.incomplete`

```
1 {
2   "id": "evt_abc123",
3   "type": "response.incomplete",
4   "created_at": 1719168000,
5   "data": {
6     "id": "resp_abc123"
7   }
8 }
```

batch.completed

Sent when a batch API request has been completed.

created_at integer
The Unix timestamp (in seconds) of when the batch API request was completed.

data object
Event data payload.
▼ Show properties

id string
The unique ID of the event.

object string
The object of the event. Always `event` .

type string
The type of the event. Always `batch.completed` .

OBJECT batch.completed

```
1 {
2   "id": "evt_abc123",
3   "type": "batch.completed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "batch_abc123"
7   }
8 }
```

batch.cancelled

Sent when a batch API request has been cancelled.

created_at integer
The Unix timestamp (in seconds) of when the batch API request was cancelled.

data object
Event data payload.
▼ Show properties

id string
The unique ID of the event.

object string
The object of the event. Always `event` .

type string
The type of the event. Always `batch.cancelled` .

OBJECT batch.cancelled

```
1 {
2   "id": "evt_abc123",
3   "type": "batch.cancelled",
4   "created_at": 1719168000,
5   "data": {
6     "id": "batch_abc123"
7   }
8 }
```

batch.expired

Sent when a batch API request has expired.

created_at integer

The Unix timestamp (in seconds) of when the batch API request expired.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event`.

type string

The type of the event. Always `batch.expired`.

OBJECT batch.expired



```
1 {
2   "id": "evt_abc123",
3   "type": "batch.expired",
4   "created_at": 1719168000,
5   "data": {
6     "id": "batch_abc123"
7   }
8 }
```

batch.failed

Sent when a batch API request has failed.

created_at integer

The Unix timestamp (in seconds) of when the batch API request failed.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `batch.failed` .

OBJECT `batch.failed`



```
1 {
2   "id": "evt_abc123",
3   "type": "batch.failed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "batch_abc123"
7   }
8 }
```

`fine_tuning.job.succeeded`

Sent when a fine-tuning job has succeeded.

created_at integer

The Unix timestamp (in seconds) of when the fine-tuning job succeeded.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `fine_tuning.job.succeeded` .

OBJECT `fine_tuning.job.succeeded`



```
1 {
2   "id": "evt_abc123",
3   "type": "fine_tuning.job.succeeded",
4   "created_at": 1719168000,
5   "data": {
6     "id": "ftjob_abc123"
7   }
8 }
```

fine_tuning.job.failed

Sent when a fine-tuning job has failed.

created_at integer
The Unix timestamp (in seconds) of when the fine-tuning job failed.

data object
Event data payload.
▼ Show properties

id string
The unique ID of the event.

object string
The object of the event. Always `event` .

type string
The type of the event. Always `fine_tuning.job.failed` .

OBJECT fine_tuning.job.failed

```
1 {
2   "id": "evt_abc123",
3   "type": "fine_tuning.job.failed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "ftjob_abc123"
7   }
8 }
```

fine_tuning.job.cancelled

Sent when a fine-tuning job has been cancelled.

created_at integer
The Unix timestamp (in seconds) of when the fine-tuning job was cancelled.

data object
Event data payload.
▼ Show properties

id string
The unique ID of the event.

object string
The object of the event. Always `event` .

type string
The type of the event. Always `fine_tuning.job.cancelled` .

OBJECT fine_tuning.job.cancelled

```
1 {
2   "id": "evt_abc123",
3   "type": "fine_tuning.job.cancelled",
4   "created_at": 1719168000,
5   "data": {
6     "id": "ftjob_abc123"
7   }
8 }
```

eval.run.succeeded

Sent when an eval run has succeeded.

created_at integer

The Unix timestamp (in seconds) of when the eval run succeeded.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `eval.run.succeeded` .

OBJECT `eval.run.succeeded`



```
1 {
2   "id": "evt_abc123",
3   "type": "eval.run.succeeded",
4   "created_at": 1719168000,
5   "data": {
6     "id": "evalrun_abc123"
7   }
8 }
```

eval.run.failed

Sent when an eval run has failed.

created_at integer

The Unix timestamp (in seconds) of when the eval run failed.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `eval.run.failed` .

OBJECT `eval.run.failed`



```
1 {
2   "id": "evt_abc123",
3   "type": "eval.run.failed",
4   "created_at": 1719168000,
5   "data": {
6     "id": "evalrun_abc123"
7   }
8 }
```

eval.run.canceled

Sent when an eval run has been canceled.

created_at integer

The Unix timestamp (in seconds) of when the eval run was canceled.

data object

Event data payload.

▼ Show properties

id string

The unique ID of the event.

object string

The object of the event. Always `event` .

type string

The type of the event. Always `eval.run.canceled` .

OBJECT `eval.run.canceled`



```
1 {
2   "id": "evt_abc123",
3   "type": "eval.run.canceled",
4   "created_at": 1719168000,
5   "data": {
6     "id": "evalrun_abc123"
7   }
8 }
```